

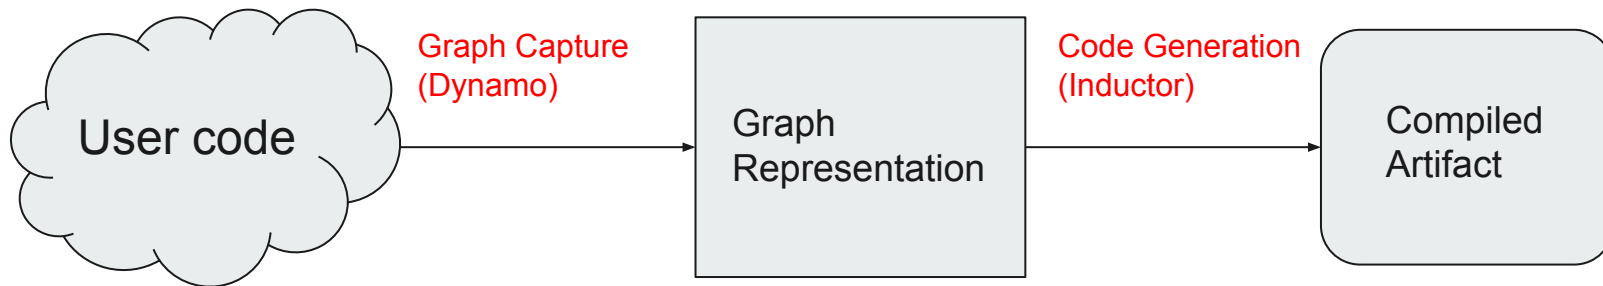


AOTAutograd Overview

Background

Overall compilation flow

1. Capture user code into a graph representation
2. Compile the graph representation into efficient code





Plan

- Problem Statement (what is AOTAutograd solving)
- Automatic Differentiation (background)
- PyTorch's C++ autograd engine (eager example)
- AOTAutograd: How we handle tracing the autograd engine in torch.compile
- Other things AOTAutograd does: functionalization
- Dynamo vs. AOTAutograd tracing: differences



Problem Statement

- Torch.compile should support training



Problem Statement

- Torch.compile should support training
- Training support requires capturing + compiling a backward graph
- **Autograd is implemented in C++ (PyTorch internals)**
 - Dynamo cannot trace into PyTorch's autograd code

Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



Background: Autodiff



Background: Autodiff

Steps when training a neural network:

1. Forward propagation
 - the network makes its best guess about the correct output. It runs the input data through each of its functions to make this guess
2. Loss function
 - Compute a scalar “loss”, dictating how far off the network was from the expected output
3. Backward propagation
 - Compute gradients which inform us of the direction in which we should move the network's weights to minimize the loss
4. Optimizer step:
 - Update the network weights given the computed gradients

```
out = model(input)  
loss = loss_fn(out, expected_out)  
out.sum().backward()  
optimizer.step()
```



Background: Autodiff

Gradient compute: derived automatically
from the forward.

The user does not write python code
corresponding to their backward.

```
out = model(input)
loss = loss_fn(out, expected_out)
out.sum().backward()
optimizer.step()
```


Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



PyTorch: Autograd



PyTorch: Autograd

- PyTorch's autograd engine is tape-based
 - Calling operators on tensors will record their backward formulas into a "tape"
 - Every operator has a mapping to its derivative formula
 - $\sin(x) \rightarrow \cos(x)$
- Invoking `.backward()` will:
 - Execute each operator in the backward tape
 - Populate gradients into the `.grad` field

```
>>> x = torch.ones(4, requires_grad=True)
>>> out = x.sin()
>>> print(out.grad_fn)
<SinBackward0 object at 0x1051b0730>
>>> out.sum().backward()
>>> print(x.grad)
tensor([0.5403, 0.5403, 0.5403, 0.5403])
```

PyTorch: Autograd

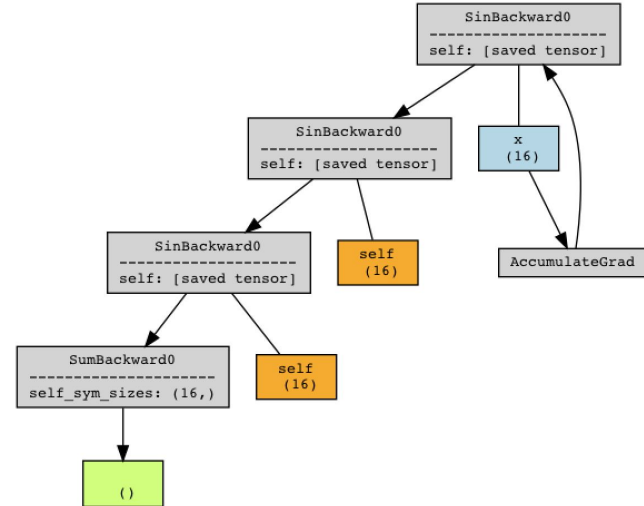
We can visualize the backward graph created by the autograd engine!

```
import torch
from torchviz import make_dot
```

```
def f(x):
    return x.sin().sin().sin()
```

```
param = torch.randn(16, requires_grad=True)
out_expected = torch.zeros(16)
out = f(x)
loss = (out - out_expected) ** 2
```

```
make_dot(
    out.sum(),
    params={'x': param},
    show_attrs=True,
    show_saved=True
).render("bw_graph", format="png")
```





AOTAutograd

Back to `torch.compile`.

Training support: we want `torch.compile` to be able to compile both the forward and the backward

- The autograd engine logic is in C++
- No bytecode for dynamo to trace
- How do we get the backward graph?

AOTAutograd

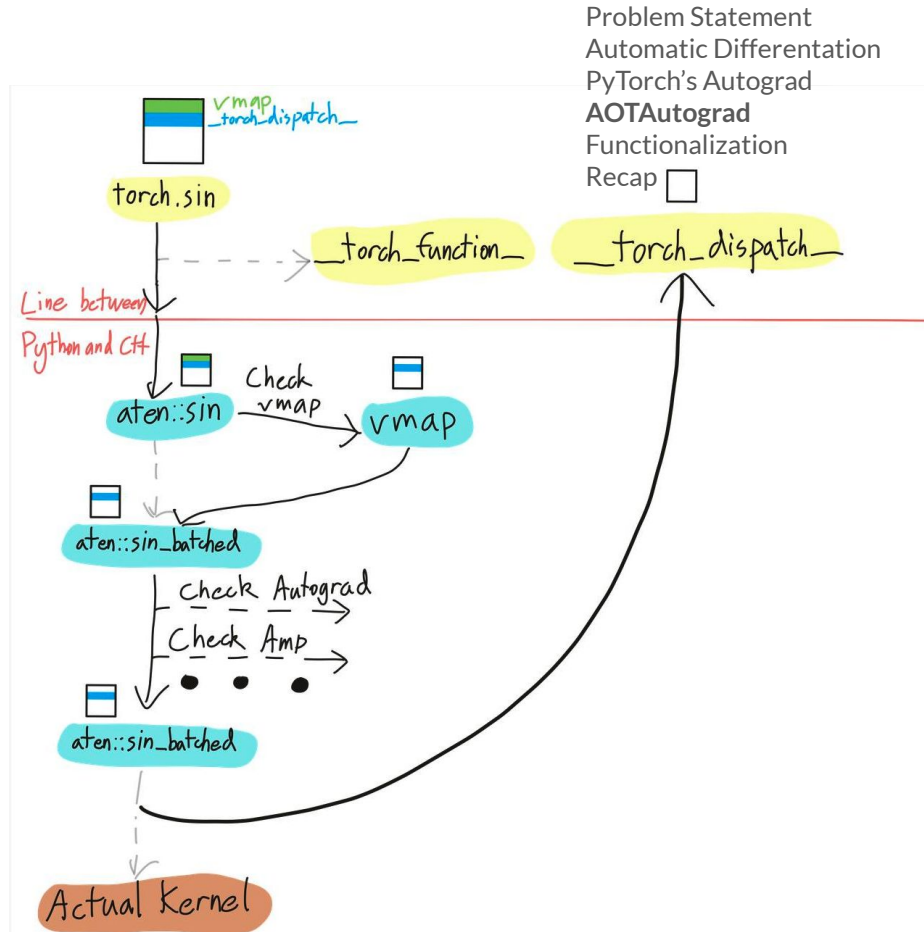
Back to torch.compile.

Training support: we want torch.compile to be able to compile both the forward and the backward

- The autograd engine logic is in C++
- No bytecode for dynamo to trace
- How do we get the backward graph?

`_torch_dispatch_`: a hook back into python, right before every operator is executed

When autograd executes its backward operations, intercept and record each operator into an FX graph



Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



AOTAutograd: Example

```
@torch.compile                                     User code
def f(x):
    return x.sin().sin().sin()
```

AOTAutograd: Example

Step 1:

- Dynamo traces the python bytecode from the user
- Traces out an FX graph containing the user's torch operations

```
===== __compiled_fn_0 =====
<eval_with_key>.0 class GraphModule(torch.nn.Module):
  def forward(self, L_x_ : torch.Tensor):
    L_x_ = L_x_

    # File: /Users/hirsheybar/tmp.py:25, code: return x.sin().sin().sin()
    sin = L_x_.sin(); L_x_ = None
    sin_1 = sin.sin(); sin = None
    sin_2 = sin_1.sin(); sin_1 = None
    return (sin_2,)
```

1 graph output

<pre>@torch.compile def f(x): return x.sin().sin().sin()</pre>	User code
--	-----------



AOTAutograd: Example

Step 2:

- AOTAutograd takes the “forward” graph from Dynamo
- Traces through the autograd engine
- Generates separate graphs for the forward and backward

```
@torch.compile                                     User code
def f(x):
    return x.sin().sin().sin()
```


AOTAutograd: Example

Step 2:

- AOTAutograd takes the “forward” graph from Dynamo
- Traces through the autograd engine
- Generates separate graphs for the forward and backward

Forward graph

2 graph outputs:

- The user output (result of sin)
- Saved activation, used in the backward pass

```
===== Forward graph 0 =====  
<eval_with_key>.35 class GraphModule(torch.nn.Module):  
    def forward(self, primals_1: "f32[16]"):   
        sin: "f32[16]" = torch.ops.aten.sin.default(primals_1)  
        sin_1: "f32[16]" = torch.ops.aten.sin.default(sin)  
        sin_2: "f32[16]" = torch.ops.aten.sin.default(sin_1)  
        return [sin_2, primals_1]
```

```
@torch.compile User code  
def f(x):  
    return x.sin().sin().sin()
```

AOTAutograd: Example

Step 2:

- AOTAutograd takes the “forward” graph from Dynamo
- Traces through the autograd engine
- Generates separate graphs for the forward and backward

Backward graph

1 graph output:

- Gradient of output w.r.t. x

```
===== Backward graph 0 =====
<eval_with_key>.36 class GraphModule(torch.nn.Module):
  def forward(self, primals_1: "f32[16]", tangents_1: "f32[16]"):
    sin: "f32[16]" = torch.ops.aten.sin.default(primals_1)
    sin_1: "f32[16]" = torch.ops.aten.sin.default(sin)
    cos: "f32[16]" = torch.ops.aten.cos.default(sin_1)
    mul: "f32[16]" = torch.ops.aten.mul.Tensor(tangents_1, cos)
    cos_1: "f32[16]" = torch.ops.aten.cos.default(sin)
    mul_1: "f32[16]" = torch.ops.aten.mul.Tensor(mul, cos_1)
    cos_2: "f32[16]" = torch.ops.aten.cos.default(primals_1)
    mul_2: "f32[16]" = torch.ops.aten.mul.Tensor(mul_1, cos_2)
    return [mul_2]
```

@torch.compile

User code

```
def f(x):
    return x.sin().sin().sin()
```



What happens at runtime

```
def f(x):  
    return x.sin().sin().sin()  
  
x = torch.randn(16, requires_grad=True)  
out = f(x)  
print(out.grad_fn)
```

User code

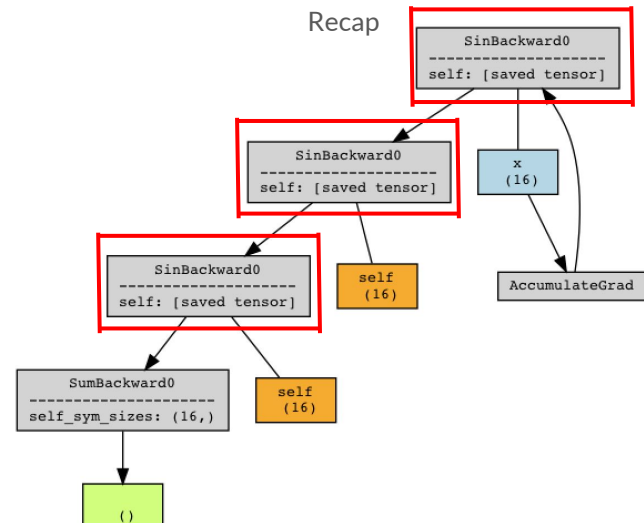
What happens at runtime

In eager mode

- 3 ops in backward
- Every sin() saves its input for backward

<SinBackward0 object at 0x150b22710>

Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



```
def f(x):  
    return x.sin().sin().sin()  
  
x = torch.randn(16, requires_grad=True)  
out = f(x)  
print(out.grad_fn)
```

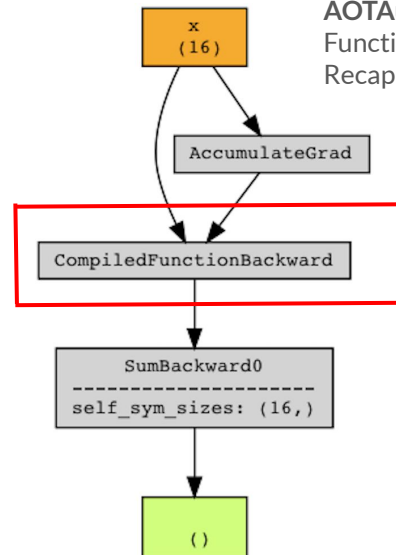
User code

What happens at runtime

With torch.compile

- 1 op in backward: “Compiled backward”
- No need to save all 3 tensors for backward

Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



```
def f(x):  
    return x.sin().sin().sin()  # User code  
  
x = torch.randn(16, requires_grad=True)  
out = f(x)  
print(out.grad_fn)
```

Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



What else does AOTAutograd handle



What else does AOTAutograd handle

AOTAutograd traces “framework” code that lives in C++

- Motivating use case: autograd engine (for training support)
- But many other functionalities in PyTorch are implemented in C++ framework code
 - (for eager performance)
- AOTAutograd traces through these too
 - AMP (automatic mixed precision)
 - Functorch transforms (vmap/grad)
 - Tensor subclasses (user-land extension point)
 - Operator decompositions
 - **Functionalization (remove mutations from a program)**

Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



Functionalization



Functionalization

- Another transform that lives in C++
- Removes mutations from a program
 - Compilers prefer a functional graph
- AOTAutograd traces through it too

```
@torch.compile(backend="aot_eager")
def f(x):
    a = x.add(1)
    a.add_(2)
    return a.add(3)

x = torch.ones(4, 4)
out = f(x)
```

Functionalization

- Another transform that lives in C++
- Removes mutations from a program
 - Compilers prefer a functional graph
- AOTAutograd traces through it too

Dynamo graph: has mutation

```
def forward(self, l_x_ : torch.Tensor):  
    # File: /Users/hirsheybar/tmp2.py:5, code: a = x.add(1)  
    a = l_x_.add(1)  
  
    # File: /Users/hirsheybar/tmp2.py:6, code: a.add_(2)  
    add_ = a.add_(2)  
  
    # File: /Users/hirsheybar/tmp2.py:7, code: return a.add(3)  
    add_1 = a.add(3)  
    return (add_1,)
```

```
@torch.compile(backend="aot_eager")  
def f(x):  
    a = x.add(1)  
    a.add_(2)  
    return a.add(3)  
  
x = torch.ones(4, 4)  
out = f(x)
```

Functionalization

- Another transform that lives in C++
- Removes mutations from a program
 - Compilers prefer a functional graph
- AOTAutograd traces through it too

AOTAutograd graph: functional

```
def forward(self, arg0_1: "f32[4, 4]"):
    # File: /Users/hirsheybar/tmp2.py:5, code: a = x.add(1)
    add: "f32[4, 4]" = torch.ops.aten.add.Tensor(arg0_1, 1)

    # File: /Users/hirsheybar/tmp2.py:6, code: a.add_(2)
    add_1: "f32[4, 4]" = torch.ops.aten.add.Tensor(add, 2)

    # File: /Users/hirsheybar/tmp2.py:7, code: return a.add(3)
    add_2: "f32[4, 4]" = torch.ops.aten.add.Tensor(add_1, 3)
    return (add_2,)
```

```
@torch.compile(backend="aot_eager")
def f(x):
    a = x.add(1)
    a.add_(2)
    return a.add(3)

x = torch.ones(4, 4)
out = f(x)
```

Problem Statement
Automatic Differentiation
PyTorch's Autograd
AOTAutograd
Functionalization
Recap



Back to the overview



Back to the overview

Dynamo:

- traces the user's python bytecode, puts all torch.* operators into a graph
- Compiling this graph is hard! Mutation, aliasing, backward is implicitly defined

AOTAutograd

- Generates a “simpler” graph, given the graph from dynamo
- Does so by tracing the PyTorch “framework” code that lives in C++

User code

(good for Dynamo)

- All in Python
- Can do arbitrarily crazy things
 - Lots of global state
 - 3rd party libs
- Mix of “tensor compute” and python side effects

(PyTorch) Framework code

(good for AOTAutograd)

- Mostly in C++
- Does not do arbitrarily crazy things (easy to trace)
 - We control it (can tweak make it tracer-friendly)
 - No 3rd party libs
- 100% tensor compute