

TorchDynamo Deep Dive

Agenda

- Motivation
- TorchDynamo bytecode analysis
- TorchDynamo components
- Practical session

TorchDynamo: Motivation

The Great ML Framework Debate

Eager Mode

- Preferred by users
- Easier to use programming model
- Easy to debug
- PyTorch is a primarily an eager mode framework

Graph Mode

- Preferred by backends and framework builders
- Easier to optimize with a compiler
- Easier to do automated transformations

PyTorch's Many Attempts at Graph Modes

torch.jit.trace

- Record + replay
- Unsound
- Can give incorrect results because it ignores Python part of program

torch.jit.script

- AOT parses Python into graph format
- Only works on ~45% of real world models
- High effort to “TorchScript” models

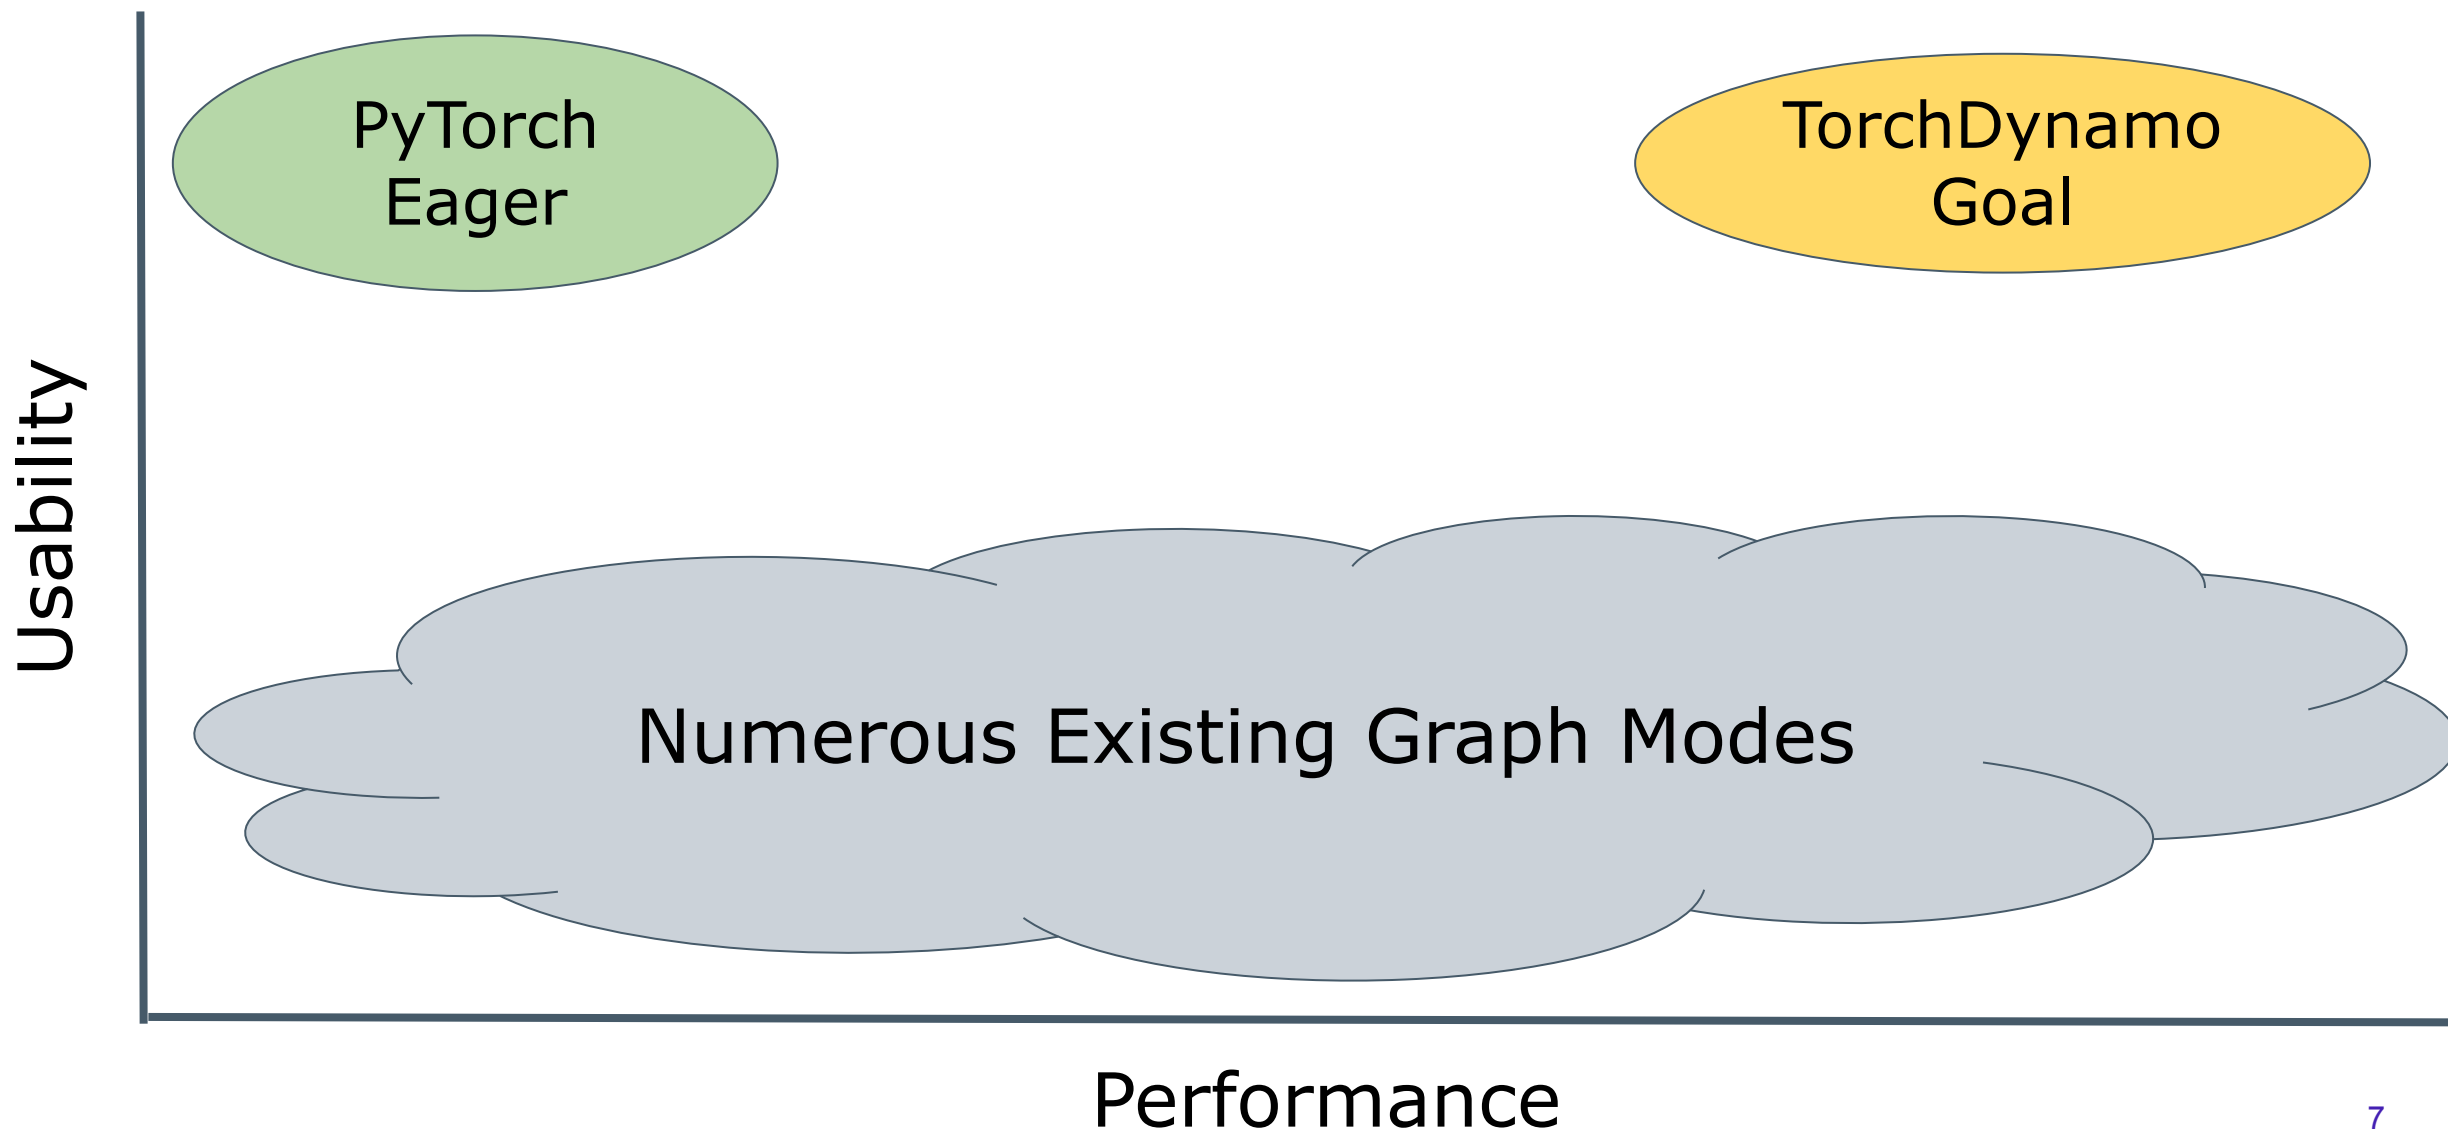
PyTorch Models Are Not Static Graphs

PyTorch users write models where program graphs are impossible

- Convert tensors native Python types (`x.item()`, `x.tolist()`, `int(x)`, etc)
- Use other frameworks (numpy/xarray/etc) for part of their model
- Data dependent Python control flow or other dynamism
- Exceptions, closures, generators, classes, etc

These violate the assumptions of most graph mode backends

PyTorch Usability/Performance Tradeoff



7k+

Crawled GitHub Models


TorchDynamo

Python-level JIT compiler to make unmodified PyTorch programs run faster


- Sprinkles graphs in eager mode
- Fallback to Python for hard-to-accelerate (non-graph) code

TorchDynamo: How does it work?

Default Python Behavior



TorchDynamo Behavior



Toy Example


```
def fn(x):  
    a = torch.sin(x)  
    b = torch.cos(x)  
    return a * b  
  
x = torch.randn(10)  
opt_fn = torch._dynamo.optimize("eager")(fn)  
opt_fn(x)
```

When `opt_fn` is called, TorchDynamo takes control:


- `custom_eval_frame(PyFrameObject* frame)`
 - `frame->f_locals`
 - `{"x": tensor([...]), "y": tensor([...])}`
 - `frame->f_globals`
 - `{"torch": ..., ...}`
 - `frame->f_code`
 - Bytecode
 - ...
 - ...

Memory Offset	Instruction	Argument Raw (Decoded)
0	LOAD_GLOBAL	0 (torch)
2	LOAD_METHOD	1 (sin)
4	LOAD_FAST	0 (x)
6	CALL_METHOD	1
8	STORE_FAST	2 (a)
10	LOAD_GLOBAL	0 (torch)
12	LOAD_METHOD	2 (cos)
14	LOAD_FAST	0 (x)
16	CALL_METHOD	1
18	STORE_FAST	3 (b)
20	LOAD_FAST	2 (a)
22	LOAD_FAST	3 (b)
24	BINARY_MULTIPLY	
26	RETURN_VALUE	

Default Python Behavior



TorchDynamo Behavior



TorchDynamo: Bytecode Analysis and Graph Extraction

```
0 LOAD_GLOBAL          0 (torch)
2 LOAD_METHOD          1 (sin)
4 LOAD_FAST            0 (x)
6 CALL_METHOD          1
8 STORE_FAST           2 (a)

10 LOAD_GLOBAL          0 (torch)
12 LOAD_METHOD          2 (cos)
14 LOAD_FAST            0 (x)
16 CALL_METHOD          1
18 STORE_FAST           3 (b)

20 LOAD_FAST            2 (a)
22 LOAD_FAST            3 (b)
24 BINARY_MULTIPLY
26 RETURN_VALUE
```

Original Python Bytecode



```
def forward(self, x : torch.Tensor):
    sin = torch.sin(x)
    cos = torch.cos(x); x = None
    mul = sin * cos; sin = cos = None
    return (mul,)
```

Extracted Fx Graph



```
0 LOAD_GLOBAL          3 (__compiled_fn_0)
2 LOAD_FAST            0 (x)
4 CALL_FUNCTION         1
6 UNPACK_SEQUENCE       1
8 RETURN_VALUE
```

Optimized Python Bytecode

TorchDynamo: Integration with Backend Compiler

```
0 LOAD_GLOBAL          0 (torch)
2 LOAD_METHOD          1 (sin)
4 LOAD_FAST            0 (x)
6 CALL_METHOD          1
8 STORE_FAST           2 (a)

10 LOAD_GLOBAL          0 (torch)
12 LOAD_METHOD          2 (cos)
14 LOAD_FAST            0 (x)
16 CALL_METHOD          1
18 STORE_FAST           3 (b)

20 LOAD_FAST            2 (a)
22 LOAD_FAST            3 (b)
24 BINARY_MULTIPLY
26 RETURN_VALUE
```

Original Python Bytecode

```
def forward(self, x : torch.Tensor):
    sin = torch.sin(x)
    cos = torch.cos(x); x = None
    mul = sin * cos; sin = cos = None
    return (mul,)
```

Extracted Fx Graph


```
0 LOAD_GLOBAL          3 (__compiled_fn_0)
2 LOAD_FAST            0 (x)
4 CALL_FUNCTION         1
6 UNPACK_SEQUENCE       1
8 RETURN_VALUE
```

Optimized Python Bytecode

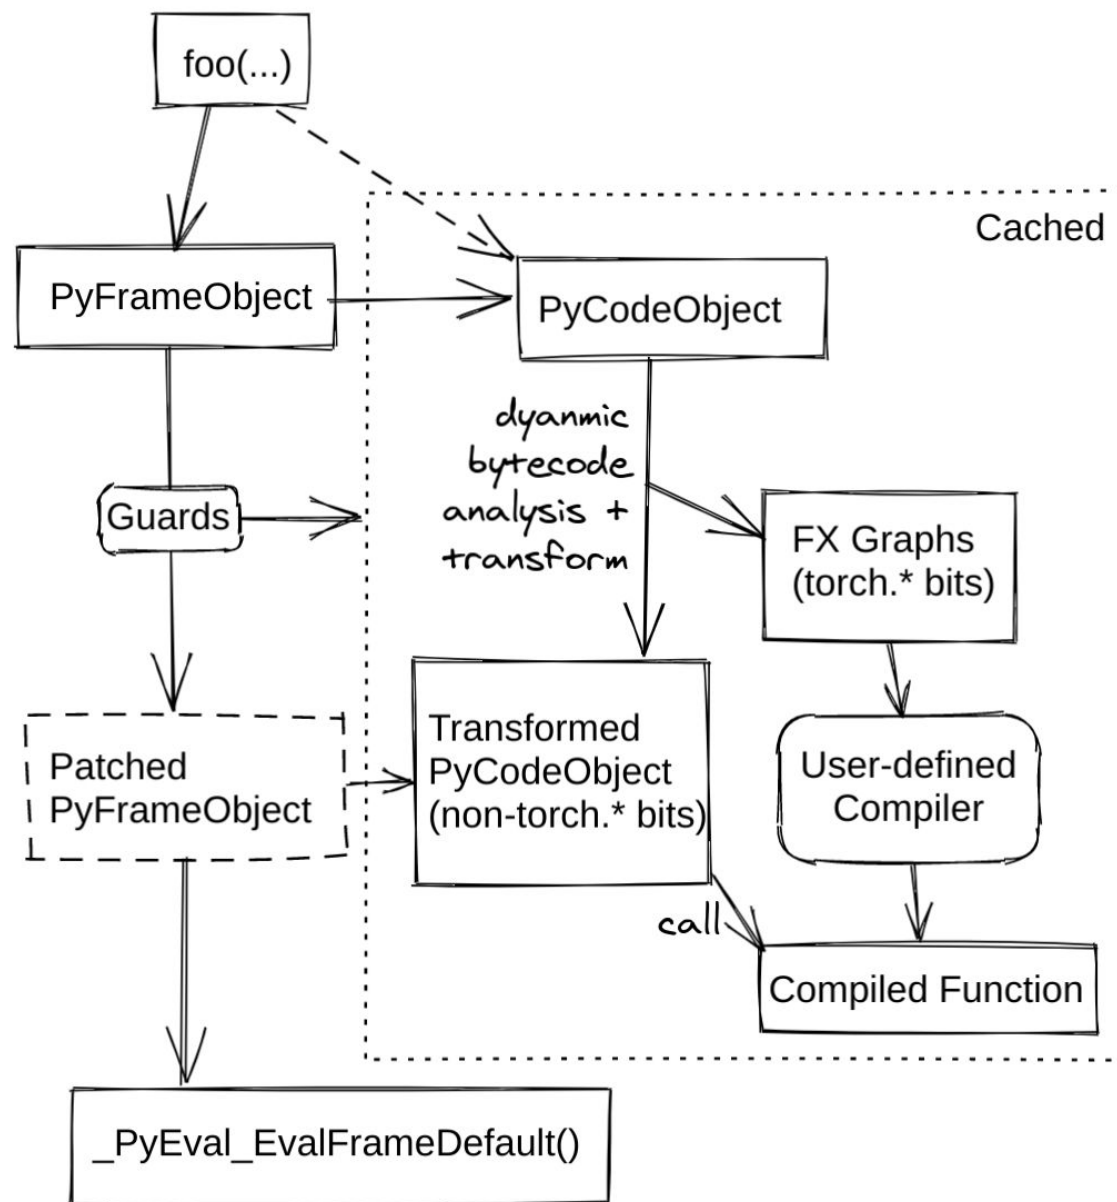
```
def my_compiler(gm, example_inputs):
    return optimized_fn(gm, example_inputs)
```

Backend compilers accept a Fx graph module and return optimized callable

Default Python Behavior



TorchDynamo Behavior



TorchDynamo Concepts

Google Collab Demo

https://colab.research.google.com/drive/19JURKGhy_L82Y-2MUc2jurJwARCPy-YL?usp=sharing