


PyTorch Export

Sound Whole-Graph Capture for PyTorch

Avik Chaudhuri (PyTorch Compiler, AI at Meta)





`torch.compile`


`torch.export`

just-in-time compiler

needs Python runtime

emits **backend-specific code**

doesn't need source changes




torch.compile	torch.export
just-in-time compiler	ahead-of-time frontend compiler
needs Python runtime	cuts dependency on Python runtime
emits backend-specific code	emits backend-agnostic IR
doesn't need source changes	compile-time errors possible

AOT (instead of JIT) compilation is often...


- **necessary**

- Python-free environments
 - e.g., on-device, serving



- **desirable**

- backend-specific optimizations
 - e.g., custom kernels, special hardware
- stability
 - e.g., online/recurring training



Graph breaks expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x):
        if x.sum() > 0:
            return x.sin()
        else:
            return x + 1
```

dynamic control flow

```
m = M()
```

Graph breaks expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x):
        if x.sum() > 0:
            return x.sin()
        else:
            return x + 1
```

```
m = M()
m(torch.randn(8))
```

```
def graph_0(x):
    sum = x.sum()
    gt = sum > 0
    return gt
```

```
def graph_1(x):
    add = x + 1
    return add
```



**interpreter
fallback**

Graph breaks considered harmful for AOT

```
class M(torch.nn.Module):  
    def forward(self, x):  
        return torch.cond(  
            x.sum() > 0,  
            lambda x: x.sin(),  
            lambda x: x + 1,  
            (x, ),  
        )
```

rewrite

```
m = M()  
torch.export(m, (torch.randn(8),))
```

```
def graph_0(x):  
    sum = x.sum()  
    gt = sum > 0  
  
    def true_fn(x):  
        sin = x.sin()  
        return sin  
  
    def false_fn(x):  
        add = x + 1  
        return add  
  
    cond = torch.ops.higher_order.cond(  
        gt, true_fn, false_fn, (x, )  
    )  
    return cond
```

Recompilations expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x, y):
        if x.shape[0] <= 1024:
            return x + y
        else:
            return x
```

static control flow

```
m = M()
```


Slogan: shapes are types!

complexity ↓	Python types (e.g., zip)	PyTorch shapes (e.g., matmul)
	<code>(List, List) -> List</code>	<code>(Tensor, Tensor) -> Tensor</code>

Slogan: shapes are types!

complexity ↓	Python types (e.g., zip)	PyTorch shapes (e.g., matmul)
	<code>(List, List) -> List</code>	<code>(Tensor, Tensor) -> Tensor</code>
	<code>(List[int], List[str]) -> List[(int, str)]</code>	<code>(Tensor[4,8], Tensor[8,16]) -> Tensor[4,16]</code>

Slogan: shapes are types!

complexity ↓	Python types (e.g., zip)	PyTorch shapes (e.g., matmul)
	<code>(List, List) -> List</code>	<code>(Tensor, Tensor) -> Tensor</code>
	<code>(List[int], List[str]) -> List[(int, str)]</code>	<code>(Tensor[4,8], Tensor[8,16]) -> Tensor[4,16]</code>
	<code>(List[X], List[Y]) -> List[(X, Y)]</code>	<code>(Tensor[a,b], Tensor[b,c]) -> Tensor[a,c]</code>

Recompilations expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x, y):
        if x.shape[0] <= 1024:
            return x + y
        else:
            return x
```

static control flow

```
m = M()
```

Recompilations expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x, y):
        if x.shape[0] <= 1024:
            return x + y
        else:
            return x
```

```
m = M()
m(torch.randn(512), torch.randn(512))
```

```
def graph_0(x: f[512], y: f[512]):
    add: f[512] = x + y
    return add
```

static shapes

Recompilations expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x, y):
        if x.shape[0] <= 1024:
            return x + y
        else:
            return x
```

```
m = M()
m(torch.randn(512), torch.randn(512))
m(torch.randn(1024), torch.randn(1024))
```

```
# s <= 1024
def graph_1(x: f[s], y: f[s]):
    add: f[s] = x + y
    return add
```

~~static shapes~~
dynamic shapes

Recompilations expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x, y):
        if x.shape[0] <= 1024:
            return x + y
        else:
            return x
```

```
m = M()
m(torch.randn(512), torch.randn(512))
m(torch.randn(1024), torch.randn(1024))
m(torch.randn(2048), torch.randn(1024))
```

```
# s > 1024
def graph_2(x: f[s], y):
    return x
```

~~static shapes~~
dynamic shapes
(new path)

Recompilations expected for JIT

```
@torch.compile()
class M(torch.nn.Module):
    def forward(self, x, y):
        if x.shape[0] <= 1024:
            return x + y
        else:
            return x
```

```
m = M()
m(torch.randn(512), torch.randn(512))
m(torch.randn(1024), torch.randn(1024))
m(torch.randn(2048), torch.randn(1024))
m(torch.randn(1024), torch.randn(2048))
```

```
# s <= 1024
def graph_1(x: f[s], y: f[s?]):
    add: f[s?] = x + y
    return add
```

static-shapes
dynamic-shapes
(new-path)
run-time exception

Recompilations considered harmful for AOT

```
class M(torch.nn.Module):  
    def forward(self, x, y):  
        if x.shape[0] <= 1024:  
            return x + y  
        else:  
            return x
```

```
m = torch.export(M(), (torch.randn(512), torch.randn(512))).module()
```

```
def graph_0(x: f[512], y: f[512]):  
    add: f[512] = x + y  
    return add
```

static shapes

Recompilations considered harmful for AOT

```
class M(torch.nn.Module):  
    def forward(self, x, y):  
        if x.shape[0] <= 1024:  
            return x + y  
        else:  
            return x
```

```
m = torch.export(M(), (torch.randn(512), torch.randn(512))).module()
```

```
m(torch.randn(1024), torch.randn(1024))
```

```
def graph_0(x: f[512], y: f[512]):  
    add: f[512] = x + y  
    return add
```

static shapes

**assertion failure
(shapes not dynamic)**

Recompilations considered harmful for AOT

```
class M(torch.nn.Module):  
    def forward(self, x, y):  
        if x.shape[0] <= 1024:  
            return x + y  
        else:  
            return x
```

```
args = (torch.randn(512), torch.randn(512))  
s = torch.export.Dim("s", max=1024)  
m = torch.export(M(), args, dynamic_shapes=((s), (s))).module()  
  
m(torch.randn(1024), torch.randn(1024))
```

```
# s <= 1024  
def graph_1(x: f[s], y: f[s]):  
    add: f[s] = x + y  
    return add
```

dynamic shapes

ok

Recompilations considered harmful for AOT

```
class M(torch.nn.Module):  
    def forward(self, x, y):  
        if x.shape[0] <= 1024:  
            return x + y  
        else:  
            return x
```

```
args = (torch.randn(512), torch.randn(512))  
s = torch.export.Dim("s", max=1024)  
m = torch.export(M(), args, dynamic_shapes=((s), (s))).module()
```

```
m(torch.randn(1024), torch.randn(1024))  
m(torch.randn(2048), torch.randn(1024))  
m(torch.randn(1024), torch.randn(2048))
```


```
# s <= 1024  
def graph_1(x: f[s], y: f[s]):  
    add: f[s] = x + y  
    return add
```

dynamic shapes

ok


**assertion failure
(shape constraints violated)**

Data-driven UX




Focus

- **Simplicity**
 - predictable behaviors
- **Actionable errors**
 - source locations
 - suggested fixes
- **Expressivity**
 - control flow, state
 - shape constraints




Syntactic guarantees

- **Inlining**
 - flattened graph





Syntactic guarantees

- **Inlining**
 - flattened graph
- **Functionalization**
 - lifted inputs/outputs




Syntactic guarantees

- **Inlining**
 - flattened graph
- **Functionalization**
 - lifted inputs/outputs
- **Metadata**
 - source map (unlift+unflatten)
 - shape constraints (soundness)



Syntactic guarantees

- **Inlining**
 - flattened graph
- **Functionalization**
 - lifted inputs/outputs
- **Metadata**
 - source map
 - shape constraints



Semantic guarantees

- **Behavioral correctness**
 - assume / guarantee


Syntactic guarantees

- **Inlining**
 - flattened graph
- **Functionalization**
 - lifted inputs/outputs
- **Metadata**
 - source map
 - shape constraints



Semantic guarantees

- **Behavioral correctness**
 - assume / guarantee
- **Round-tripping identities**
 - serialize / deserialize
 - unlift+unflatten / re-export



Customization [operator decomposition]

```
# “full” aten opset  
exported = ...
```

```
# “core” aten opset (standard for maximum coverage)  
exported_core = exported.run_decompositions()
```

```
# plug-in custom kernel for op  
exported_custom = exported.run_decompositions(  
    decomp_table=... # default except op  
)
```

e.g., convolution

Transformation [graph optimization]

```
# functionalized, linear FX graph  
exported = ...
```

```
# sequence of arbitrary FX transforms  
lowered = optimization_passes(exported)
```

```
# execute!  
... = backend(lowered)
```



ONNX
Torch-XLA
TensorRT
ExecuTorch

...

AOTInductor